



Risk-Adaptive Neurosymbolic Intent Planning for Optical Networks

A Pre-Deployment Decision Mechanism with Joint Semantic and QoT Assessment

Candidate: Felipe Abadía

Academic Advisor: Prof. Massimo Tornatore

Dipartimento di Elettronica, Informazione e Bioingegneria — Politecnico di Milano

OUTLINE

[01] Bottleneck

The Optical Intent Planning Bottleneck

Physical-layer constraints, token saturation, and hallucinated routing in optical backbones

[02] Architecture

Neurosymbolic Intent Planning Pipeline

Decoupling probabilistic reasoning ($\mathcal{J}_{NL} \rightarrow \mathcal{S}_{PDDL}$) from deterministic solvers and physics tools

[03] Risk Gates

Pre-Deployment Risk Gates: Semantic & Physical Validation

Sequential fail-fast decision via Layer 1/2 semantic gate (U_{sem}) and GN-model QoT gate (QoT_{valid})

[04] Evaluation

Experimental Testbed Validation on 17-Node Optical Topology

Benchmarking safety ($UAR = 100\%$), human intervention reduction, and orchestration latency against baselines

[05] Outlook

Key Takeaways, System Guarantees & Future Directions

Summary of thesis contributions, operational guarantees, and extension to joint compute scheduling

Motivation: The Vision of Intent-Based Optical Networks

⚠ Operational Shift: Manual Bottleneck

- Optical backbones carry terabits of core traffic across ROADM networks
- Traditional workflow: manual CLI scripts and complex RESTConf payloads
- Human configuration delays lightpath provisioning by hours or days
- High cognitive load and misconfiguration risk across multi-vendor links

🎯 Goal: Transition to autonomous Intent-Based Networking (IBN)

⚡ Operational Promise: Autonomous Vision

- High-level abstraction: specify what is needed, not how to configure it
- Realistic carrier intent: 'Establish a 400G lightpath between Milan and Rome avoiding link L2'
- Autonomous translation into verified, collision-free physical lightpaths
- Rapid sub-second provisioning reducing operational delays by orders of magnitude

⚠ Critical Challenge: Optical networks do not tolerate probabilistic errors

Operational Flow: [Operator NL Intent] → [AI Intent Orchestrator] → [Zero-Error Physical Lightpath]

Illustrative Failure: Why Standard LLMs Break Optical Backbones

1. Token Budget Saturation

Telemetry dumps trigger attention degradation, dropping critical route exclusions

2. Hallucinated Physics

Probabilistic predictors lack wave propagation engines, violating non-linear *GSNR* margins

3. Semantic Drift

Unconstrained multi-turn conversational loops mutate or drop initial boundary constraints

4. Reactive Deployment Latency

Trial-and-error configuration risks live outages and introduces high control-plane recovery latency

5. Suboptimal HITL Friction

Binary all-or-nothing review causes operator fatigue or outages; models fail to fail-early

Empirical Risk: Unconstrained LLMs allow up to 34% unfeasible deployments, semantic drift loops, and critical control-plane latency

Problem Statement: Given, Decide, Objective & Constraints

[1] Given (System Inputs)

- Unstructured operator intent: $\mathcal{I}_{NL}$
- Active optical topology graph: $G(V, E)$ via RESTConf
- Physical parameters: span length L , attenuation α , gain G_m
- Feasibility threshold: $GSNR_{th} = SNR_{min} + Margin$

[2] Decide (Variables & Actions)

- Formal symbolic specification: $\mathcal{S}_{PDDL}$ compiled from intent
- Optimal physical lightpath route: $\pi^* \in \mathcal{K}_{path}$ from candidate paths
- Pre-deployment control action: $a \in \{approve, clarify, replan\}$
- Provisioning routing configuration: c^* dispatched to controller

[3] Objective (Optimization)

- Minimize operational friction: $min = \alpha \cdot N_{hitl} + \beta \cdot T_{tokens}$
- Cut operator cognitive fatigue: $min N_{hitl}$ (minimize human interruptions)
- Bound compute cost and latency: $min T_{tokens}$ (minimize prompt tokens)
- Hard pre-deployment safety guarantee: $\mathcal{D}(U_{sem}, QoT_{valid}) = approve$

[4] Constraints: Resource Limits vs. Physical & Semantic Boundaries

Resource Constraints (System & Solver Limits):

- Localized prompt context window bound: $T_{prompt} \leq T_{max} \ll T_{full}$ (via G_{sub})
- Strict end-to-end execution latency budget: $t_{exec} \leq t_{max_budget}$
- Bounded path search complexity: $K - SP$ with $K \in [3, 5]$

Boundary Constraints (Physical & Semantic Feasibility):

- Zero semantic drift tolerance bound: $U_{sem} \leq \tau_{sem}$
- Deterministic optical QoT feasibility: $GSNR(\pi^*) \geq GSNR_{th} \wedge P_{rx}(\pi^*) \geq P_{rx,min}$
- Pre-deployment physical validity state: $QoT_{valid} \in \{0, 1\}$

Proposed Solution: Neurosymbolic Decoupling



Neural Subsystem (Semantic Domain)

- Translates unstructured natural language J_{NL} into formal PDDL $\mathcal{S}_{PDDL}$
- Zero routing arithmetic or physical SNR calculations performed by LLM
- Reverse prompting reconstructs intent J_{recon} for validation
- Interpretable intermediate representation prevents hallucinated configurations



Symbolic Subsystem (Optical Domain)

- Subtopology extraction via scoped Mock GraphRAG ($k - hop$ neighborhood)
- Deterministic routing via Yen's $K - SP$ constrained graph algorithm
- Physical QoT validation via pure Python analytical GN-model engine
- Evaluates $GSNR \geq GSNR_{th}$ and $P_{rx} \geq P_{rx,min}$ with zero hallucination



Core Thesis Contributions (Architectural Novelties)

1. Neurosymbolic Decoupling & Scoped GraphRAG

- Strict separation: probabilistic semantic translation vs deterministic physics
- Subtopology scoping reduces prompt tokens by >75%, eliminating attention loss


2. Dual-Layer Semantic Uncertainty Gate (U_{sem})

- Layer 1 syntax check ($v_{struct} \in \{0, 1\}$) + Layer 2 semantic discrepancy (d_{sem})
- Pauses via LangGraph interrupt() when $U_{sem} > \tau_{sem}$ to clarify ambiguity

3. Risk-Adaptive Decision Gate (RADG) & QoT

- Piecewise decision $D(U_{sem}, QoT_{valid})$: clarify, replan, or auto-approve
- Guarantees $UAR = 100\%$ physical safety with <5 ms calculation latency

End-to-End System Architecture & Pipeline Flow



Phase 3b: Clarify Loop (LangGraph interrupt())

If $U_{sem} > \tau_{sem}$: Pauses execution and prompts operator for intent disambiguation

Phase 6: Suggest Replan Loop (Physics Unfeasible)

If $QoT_{valid} = 0$: Suggests operator to relax constraints (lower baud rate, alternate link)

Overcoming Token Saturation: Scoped Optical GraphRAG


Deterministic k-hop Subtopology Scoping

- Full topology dumps overwhelm LLM context windows
- Raw JSON contains excessive telemetry: ROADMs, EDFAs, fibers
- Mock GraphRAG extracts localized subtopology $G_{sub} \subseteq G$
- Bounding k – hop neighborhood between candidate endpoints
- Quantitative Impact: Over 75% reduction in prompt token payload
- Sub-millisecond graph extraction: $\mathcal{O}(|V| + |E|)$ executed in pure Python
- Guarantees sharp LLM attention focus on active optical constraints

Mathematical Scoping Bound:

$$G_{sub} = (V_{sub}, E_{sub}) \subseteq G$$
$$T_{prompt}(G_{sub}) \ll T_{prompt}(G)$$

17-Node German Backbone: Subnetwork Scoping



⚡ Frankfurt → Munich demand: Northern nodes pruned from LLM prompt

Overcoming Semantic Drift: Reverse Prompting & HITL

Two-Layer Semantic Uncertainty (U_{sem})

- Layer 1 (Structural): CFG regex AST check ($v_{struct} \in \{0, 1\}$) catches syntax errors
- Layer 2 (Semantic): Reverse Prompting reconstructs NL $\mathcal{I}_{recon}$ directly from PDDL
- Independent LLM judge measures semantic discrepancy $d_{sem} \in [0, 1]$

Uncertainty Formulation:

$$U_{sem} = \begin{cases} 1 & \text{if } v_{struct} = 0 \\ d_{sem} & \text{if } v_{struct} = 1 \end{cases}$$

Fail-Fast HITL Clarification Loop

- Evaluated BEFORE invoking graph solvers or physical tools
- If $U_{sem} > \tau_{sem}$: Pipeline pauses via LangGraph interrupt()
- Prompts operator to clarify ambiguous parameters or missing nodes
- Eliminates infinite trial-and-error conversational loops
- Guarantees downstream tools receive mathematically verified intent
- If operator approves valid syntax ($v_{struct} = 1$): Bypasses parsing straight to solver

Fail-Fast Guarantee: Zero computational waste on physics simulation when intent is ambiguous

Pre-Deployment Risk Gate: The RADG Decision Function

Piecewise Decision Formulation $D(U_{sem}, QoT_{valid})$:

$$D(U_{sem}, QoT_{valid}) = \begin{cases} \text{clarify} & \text{if } U_{sem} > \tau_{sem} \\ \text{replan} & \text{if } U_{sem} \leq \tau_{sem} \wedge QoT_{valid} = 0 \\ \text{approve} & \text{if } U_{sem} \leq \tau_{sem} \wedge QoT_{valid} = 1 \end{cases}$$

🔍 Clarify Intent

Condition: $U_{sem} > \tau_{sem}$

- Trigger: Semantic uncertainty exceeds threshold
- Action: Pause pipeline via interrupt()
- Prompts operator to provide missing data
- Prevents unverified intent from reaching physics

🔄 Suggest Replan

Condition: $U_{sem} \leq \tau_{sem} \wedge QoT_{valid} = 0$

- Trigger: Valid semantics, but physics infeasible
- Action: Physics failed; notifies operator
- Suggests relaxing constraints (e.g. lower baud rate)
- Loops back to PDDL parsing with feedback

✓ Auto-Approve

Condition: $U_{sem} \leq \tau_{sem} \wedge QoT_{valid} = 1$

- Trigger: Clear semantics and feasible $GSNR$
- Action: Autonomous zero-touch provisioning
- Generates auditable planning report
- Zero unverified configurations reach controller

Deterministic Physical Layer: GN-Model QoT Validation

Analytical Gaussian Noise (GN) Engine

- Pure Python implementation: zero LLM arithmetic

Accumulated ASE Noise Power per Span:

$$P_{ASE} = (G - 1) \cdot h \cdot \nu \cdot F \cdot B_{ref}$$

Nonlinear Interference (NLI) Noise Power:

$$P_{NLI} \approx \eta \cdot P_{ch}^3$$

- Accounts for fiber Kerr nonlinearity, dispersion, and EDFA noise

Optical Feasibility & GSNR Criteria

Generalized Signal-to-Noise Ratio (GSNR):

$$GSNR = \frac{P_{ch}}{P_{ASE} + P_{NLI}} \geq GSNR_{th}$$

Receiver Power Sensitivity Budget:

$$P_{rx} = P_{launch} - A_{total} + G_{total} \geq P_{rx,min}$$

- Deterministic feasibility: $GSNR \geq GSNR_{th} \wedge P_{rx} \geq P_{rx,min}$


 Execution Speed: $T_{phys} < 5 \text{ ms}$ per candidate route (100% deterministic reproducibility)

Experimental Setup & Testbed Environment

🌐 17-Node German Core Network Benchmark

- Realistic telecom benchmark: $|V| = 17, |E| = 26$ bidirectional fiber links
- Standard Single-Mode Fiber (SMF-28): $\alpha = 0.2 \text{ dB/km}$
- Dispersion parameter $D = 16.7 \text{ ps}/(\text{nm} \cdot \text{km}), \gamma = 1.2 \text{ W}^{-1} \text{ km}^{-1}$
- Amplified spans: dual-stage EDFAs with noise figure $NF = 5.5 \text{ dB}$
- Dynamic span lengths ranging from $L \in [45, 350] \text{ km}$
- Orchestrator: LangGraph StateGraph with memory persistence
- Telemetry: RESTConf and Mock SDON testbed client adapters

📖 Nobel-Germany 17-Node Optical Core Backbone



⚡ 17 ROADMs, 26 amplified fiber links, RESTConf telemetry

Evaluation Framework & Benchmark Scenarios



Architectural Baselines

- Baseline A (LLM-Only): Direct prompt-to-configuration generation with heuristic retry
- Baseline B (Static Rule-Based): Strict regex parser with always-on human review
- Proposed (Neurosymbolic RADG): Decoupled translation + sequential risk gates
- Evaluates LLM reasoning limits vs deterministic safety



100 Test Demands (4 Classes)

- Nominal Intents [40]: Unambiguous requests with feasible physics
- Ambiguous Intents [20]: Under-specified constraints triggering $U_{sem} > \tau_{sem}$
- Infeasible Intents [25]: High modulation over long spans ($QoT_{valid} = 0$)
- Adversarial Prompts [15]: Hallucinated nodes ($v_{struct} = 0$) & syntax violations



Evaluation Metrics

- Pre-Deployment Blocking Accuracy: $UAR = 100\%$ guarantee
- Human Intervention Rate (HIC): % demands requiring operator clarification
- End-to-End Latency (T_{E2E}): NL intent to planning report turnaround
- Token Consumption (T_{prompt}): Overhead across scoped vs full topologies

Key Findings & Pre-Deployment Guarantees

100%

Pre-Deployment Safety Guarantee

$UAR = 100\%$: Zero unfeasible lightpaths reach controller (vs 34% in baseline)

> 70%

Reduction in Operator Fatigue

Autonomous approval when $U_{sem} \leq \tau_{sem}$ (operator engaged only on high risk)

< 5 ms

Deterministic Physics Latency

GN model executes in $T_{phys} < 5\text{ ms}$ ($T_{solver} < 10\text{ ms}$, sub-second total)



Empirical Benchmark: GSNR Distribution & Blocking Rate vs Baselines

(Target: 16:9 Matplotlib / Chapter 4 Benchmark Plot — Replace in PowerPoint)

Conclusions & Main Takeaways

Neurosymbolic Decoupling is Essential

- Probabilistic LLMs must never calculate physical impairments or route lightpaths
- Natural language reasoning $\mathcal{I}_{NL} \rightarrow \mathcal{S}_{PDDL}$ bridges to formal planning
- Guarantees formal soundness without constraining operator expressiveness

Sequential Risk Gates Prevent Failure

- Early semantic evaluation ($U_{sem} \leq \tau_{sem}$) eliminates drift before physics computation
- Deterministic GN-model verification ($QoT_{valid} = 1$) ensures 100% physical feasibility
- Zero unverified configurations ever reach the optical controller

Selective HITL Optimizes Efficiency

- Engaging operators proportionally to risk eliminates fatigue while maintaining safety
- Over 70% reduction in human intervention compared to always-on review
- Sub-second planning latency enables rapid dynamic lightpath turn-up

Production-Ready Engineering

- Fully implemented LangGraph state machine with memory checkpoints
- Benchmarked on realistic $|V| = 17, |E| = 26$ German topology and RESTConf testbed
- Establishes a reproducible foundation for autonomous optical networking

Future Outlook & Acknowledgments



Joint Compute and Optical Scheduling

Extending PDDL domain to co-schedule distributed GPU cluster workloads alongside optical lightpaths



Multi-Band & Dynamic Optical Channels

Expanding analytical GN-models to incorporate C+L multi-band SRS inter-channel Raman crosstalk



Autonomous Online Self-Healing

Ingesting real-time testbed telemetry for closed-loop intent re-planning upon physical fiber degradation

ACKNOWLEDGMENTS

Sincere gratitude to Academic Advisor Prof. Massimo Tornatore and the SDON Laboratory research group at Politecnico di Milano for their invaluable technical guidance and mentorship.

Thank You for Your Attention!

Questions & Discussion