


Computational Efficiency across All Baselines: Latency and Token Footprint

Turnaround Latency across Risk Classes (Capped at 165s; LLM-Only reaches 491s)



Token Footprint Breakdown (Useful Tokens vs. Wasted Overhead)

